

CSC 358 – Principles of Computer Networks

## Handout # 4: Interconnecting LANs; Internet Protocol (IP)

Joe Lim

Department of Mathematical and Computational  
Sciences

joe.lim@utoronto.ca

### Announcements

- **PS1 Due Friday, Feb 2<sup>nd</sup> at 11:59pm.**
  - Remember the late submission policy.
  - Remember academic integrity guidelines.
  - Submit electronically on MarkUs.
    - **NOTE. File names must be:** ps1.pdf
  - You can scan and save as a pdf file.
    - Not preferred, but acceptable.
    - Make sure scan is readable!
- Programming assignment 1
  - Have you tested Mininet?
  - Have you been able to run your VM?
  - **Start early! Start early! Start early! ...**

## Announcements – Cont'd

- This week's tutorial
  - PA1 Overview
- Discussion Board – DisCourse
  - <https://mcs.utm.utoronto.ca/forum/>


## The Story

- So far ...
  - Layers, and protocols
  - Link layer
    - Media type, encoding
    - Framing, link model
    - Error detection, correction
- This time
  - Interconnecting LANs
    - Hubs, switches, and bridges
  - The Internet Protocol

Application
Presentation
Session
Transport
Network
Data Link
Physical

## Ethernet

- Dominant wired LAN technology
- First widely used LAN technology
- Simpler, cheaper than token LANs and ATM
- Kept up with speed race: 10 Mbps – 10 Gbps




Metcalfe's  
Ethernet  
sketch

## Ethernet Uses CSMA/CD


- Carrier sense: wait for link to be idle
  - Channel idle: start transmitting
  - Channel busy: wait until idle
- Collision detection: listen while transmitting
  - No collision: transmission is complete
  - Collision: abort transmission, and send jam signal
- Random access: exponential back-off
  - After collision, wait a random time before trying again
  - After  $m^{\text{th}}$  collision, choose  $K$  randomly from  $\{0, \dots, 2^m - 1\}$
  - ... and wait for  $K \cdot 512$  bit times before trying again

## Limitations on Ethernet Length



- Latency depends on physical length of link
  - Time to propagate a packet from one end to the other
- Suppose A sends a packet at time  $t$ 
  - And B sees an idle line at a time just before  $t+d$
  - ... so B happily starts transmitting a packet
- B detects a collision, and sends jamming signal
  - But A doesn't see collision till  $t+2d$


## Limitations on Ethernet Length



- A needs to wait for time  $2d$  to detect collision
  - So, A should keep transmitting during this period
  - ... and keep an eye out for a possible collision
- Imposes restrictions on Ethernet
  - Maximum length of the wire: 2500 meters
  - Minimum length of the packet: 512 bits (64 bytes)

## Ethernet Frame Structure

- **Addresses:** source and destination MAC addresses
  - Adaptor passes frame to network-level protocol
    - If destination address matches the adaptor
    - Or the destination address is the broadcast address
  - Otherwise, adapter discards frame
- **Type:** indicates the higher layer protocol
  - Usually IP
  - But also Novell IPX, AppleTalk, ...
- **CRC:** cyclic redundancy check
  - Checked at receiver
  - If error is detected, the frame is simply dropped



## Unreliable, Connectionless Service

- **Connectionless**
  - No handshaking between sending and receiving adapter.
- **Unreliable**
  - Receiving adapter doesn't send ACKs or NACKs
  - Packets passed to network layer can have gaps
  - Gaps will be filled if application is using TCP
  - Otherwise, the application will see the gaps


## Shuttling Data at Different Layers

- Different devices switch different things
  - Physical layer: electrical signals (repeaters and hubs)
  - Link layer: frames (bridges and switches)
  - Network layer: packets (routers)




## Physical Layer: Repeaters

- Distance limitation in local-area networks
  - Electrical signal becomes weaker as it travels
  - Imposes a limit on the length of a LAN
- Repeaters join LANs together
  - Analog electronic device
  - Continuously monitors electrical signals on each LAN
  - Transmits an amplified copy



## Physical Layer: Hubs

- Joins multiple input lines electrically
  - Designed to hold multiple line cards
  - Do not necessarily amplify the signal
- Very similar to repeaters
  - Also operates at the physical layer




## Limitations of Repeaters and Hubs

- One large collision domain
  - Every bit is sent everywhere
  - So, aggregate throughput is limited
  - E.g., three departments each get 10 Mbps independently
  - ... and then connect via a hub and must share 10 Mbps
- Cannot support multiple LAN technologies
  - Does not buffer or interpret frames
  - So, can't interconnect between different rates or formats
  - E.g., 10 Mbps Ethernet and 100 Mbps Ethernet
- Limitations on maximum nodes and distances
  - Does not circumvent the limitations of shared media
  - E.g., still cannot go beyond 2500 meters on Ethernet


## Link Layer: Bridges

- Connects two or more LANs at the link layer
  - Extracts destination address from the frame
  - Looks up the destination in a table
  - Forwards the frame to the appropriate LAN segment
- Each segment is its own collision domain



## Link Layer: Switches

- Typically connects individual computers
  - A switch is essentially the same as a bridge
  - ... though typically used to connect hosts, not LANs
- Like bridges, support concurrent communication
  - Host A can talk to C, while B talks to D






## Dedicated Access and Full Duplex

- **Dedicated access**
  - Host has direct connection to the switch
  - ... rather than a shared LAN connection
- **Full duplex**
  - Each connection can send in both directions
  - Host sending to switch, and host receiving from switch
  - E.g., in 10BaseT and 100Base T
- **Completely avoids collisions**
  - Each connection is a bidirectional point-to-point link
  - No need for carrier sense, collision detection, and so on

## Bridges/Switches: Traffic Isolation

- **Switch breaks subnet into LAN segments**
- **Switch filters packets**
  - Frame only forwarded to the necessary segments
  - Segments become separate collision domains



## Advantages Over Hubs/Repeaters


- Only forwards frames as needed
  - Filters frames to avoid unnecessary load on segments
  - Sends frames only to segments that need to see them
- Extends the geographic span of the network
  - Separate collision domains allow longer distances
- Improves privacy by limiting scope of frames
  - Hosts can “snoop” the traffic traversing their segment
  - ... but not all the rest of the traffic
- Applies carrier sense and collision detection
  - Does not transmit when the link is busy
  - Applies exponential back-off after a collision
- Joins segments using different technologies

## Disadvantages Over Hubs/Repeaters

- Delay in forwarding frames
  - Bridge/switch must receive and parse the frame
  - ... and perform a look-up to decide where to forward
  - Storing and forwarding the packet introduces delay
  - Solution: cut-through switching
- Need to learn where to forward frames
  - Bridge/switch needs to construct a forwarding table
  - Ideally, without intervention from network administrators
  - Solution: self-learning
- Higher cost
  - More complicated devices that cost more money


## Motivation For Cut-Through Switching

- Buffering a frame takes time
  - Suppose  $L$  is the length of the frame
  - And  $R$  is the transmission rate of the links
  - Then, receiving the frame takes  $L/R$  time units
- Buffering delay can be a high fraction of total delay
  - Propagation delay is small over short distances
  - Making buffering delay a large fraction of total
  - Analogy: large group walking through NYC




## Cut-Through Switching

- Start transmitting as soon as possible
  - Inspect the frame header and do the look-up
  - If outgoing link is idle, start forwarding the frame
- Overlapping transmissions
  - Transmit the head of the packet via the outgoing link
  - ... while still receiving the tail via the incoming link
  - Analogy: different folks crossing different intersections




## Motivation For Self Learning

- Switches forward frames selectively
  - Forward frames only on segments that need them
- Switch table
  - Maps destination MAC address to outgoing interface
  - Goal: construct the switch table automatically



## Self Learning: Building the Table


- When a frame arrives
  - Inspect the source MAC address
  - Associate the address with the incoming interface
  - Store the mapping in the switch table
  - Use a time-to-live field to eventually forget the mapping



## Self Learning: Handling Misses

- When frame arrives with unfamiliar destination
  - Forward the frame out all of the interfaces
  - ... except for the one where the frame arrived
  - Hopefully, this case won't happen very often

When in  
doubt,  
shout!



## Switch Filtering/Forwarding

### When switch receives a frame:

index switch table using MAC dest address

if entry found for destination

then {

if dest on segment from which frame arrived

then drop the frame

else forward the frame on interface indicated


}

else flood

forward on all but the interface  
on which the frame arrived


## Flooding Can Lead to Loops

- Switches sometimes need to broadcast frames
  - Upon receiving a frame with an unfamiliar destination
  - Upon receiving a frame sent to the broadcast address
- Broadcasting is implemented by flooding
  - Transmitting frame out every interface
  - ... except the one where the frame arrived
- Flooding can lead to forwarding loops
  - E.g., if the network contains a cycle of switches
  - Either accidentally, or by design for higher reliability




## Solution: Spanning Trees

- Ensure the topology has no loops
  - Avoid using some of the links when flooding
  - ... to avoid forming a loop
- Spanning tree
  - Sub-graph that covers all vertices but contains no cycles
  - Links not in the spanning tree do not forward frames



## Constructing a Spanning Tree

- **Need a distributed algorithm**
  - Switches cooperate to build the spanning tree
  - ... and adapt automatically when failures occur
- **Key ingredients of the algorithm**
  - Switches need to elect a "root"
    - The switch with the smallest identifier
  - Each switch identifies if its interface is on the shortest path from the root
    - And exclude it from the tree if not
  - Messages (Y, d, X)
    - From node X
    - Claiming Y is the root
    - And the distance is d




## Steps in Spanning Tree Algorithm

- Initially, each switch thinks it is the root
  - Switch sends a message out every interface
  - ... identifying itself as the root with distance 0
  - Example: switch X announces (X, 0, X)
- Switches update their view of the root
  - Upon receiving a message, check the root id
  - If the new id is smaller, start viewing that switch as root
- Switches compute their distance from the root
  - Add 1 to the distance received from a neighbor
  - Identify interfaces not on a shortest path to the root
  - ... and exclude them from the spanning tree


### Example From Switch #4's Viewpoint

- Switch #4 thinks it is the root
  - Sends (4, 0, 4) message to 2 and 7
- Then, switch #4 hears from #2
  - Receives (2, 0, 2) message from 2
  - ... and thinks that #2 is the root
  - And realizes it is just one hop away
- Then, switch #4 hears from #7
  - Receives (2, 1, 7) from 7
  - And realizes this is a longer path
  - So, prefers its own one-hop path
  - And removes 4-7 link from the tree



### Example From Switch #4's Viewpoint

- Switch #2 hears about switch #1
  - Switch 2 hears (1, 1, 3) from 3
  - Switch 2 starts treating 1 as root
  - And sends (1, 2, 2) to neighbors
- Switch #4 hears from switch #2
  - Switch 4 starts treating 1 as root
  - And sends (1, 3, 4) to neighbors
- Switch #4 hears from switch #7
  - Switch 4 receives (1, 3, 7) from 7
  - And realizes this is a longer path
  - So, prefers its own three-hop path
  - And removes 4-7 link from the tree





## Robust Spanning Tree Algorithm

- Algorithm must react to failures
  - Failure of the root node
    - Need to elect a new root, with the next lowest identifier
  - Failure of other switches and links
    - Need to recompute the spanning tree
- Root switch continues sending messages
  - Periodically reannouncing itself as the root (1, 0, 1)
  - Other switches continue forwarding messages
- Detecting failures through timeout (soft state!)
  - Switch waits to hear from others
  - Eventually times out and claims to be the root

## Switches vs. Routers

- Advantages of switches over routers
  - Plug-and-play
  - Fast filtering and forwarding of frames
  - No pronunciation ambiguity (e.g., “router” vs. “rowter”)
- Disadvantages of switches over routers
  - Topology is restricted to a spanning tree
  - Large networks require large ARP tables
  - Broadcast storms can cause the network to collapse

### Comparing Hubs, Switches, & Routers

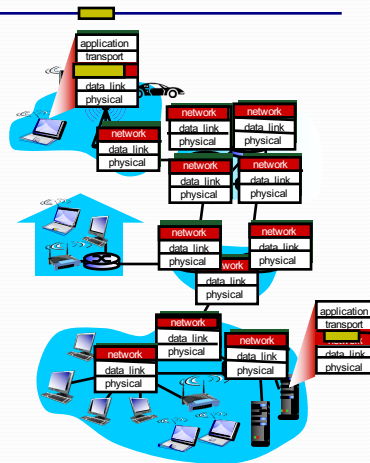
	<u>hubs</u>	<u>routers</u>	<u>switches</u>
traffic isolation	no	yes	yes
plug & play	yes	no	yes
optimal routing	no	yes	no
cut through	yes	no	yes

### Part II – The Internet Protocol (IP)

- IP: The Internet Protocol
  - Service characteristics
  - The IP Datagram format
  - IP addresses
  - Classless Inter-Domain Routing (CIDR)
  - An aside: Turning names into addresses (DNS)

## Network Layer

- transport segment from sending to receiving host
- on sending side encapsulates segments into datagrams
- on receiving side, delivers segments to transport layer
- network layer protocols in *every* host, router
- router examines header fields in all IP datagrams passing through it



## Two key Network Layer Functions


- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to dest.
  - *routing algorithms*

*analogy:*

- *routing*: process of planning trip from source to dest
- *forwarding*: process of getting through single interchange

## The Internet Protocol (IP)


### Protocol Stack



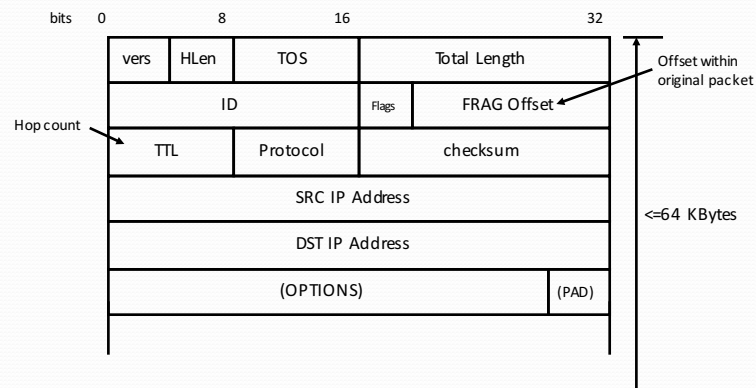
## The Internet Protocol (IP)

### Characteristics of IP

- ☒ **CONNECTIONLESS:** mis-sequencing
- ☒ **UNRELIABLE:** may drop packets...
- ☒ **BEST EFFORT:** ... but only if necessary
- ☒ **DATAGRAM:** individually routed




## The IP Datagram




## Fragmentation



**Problem:** A router may receive a packet larger than the maximum transmission unit (MTU) of the outgoing link.



**Solution:** R1 fragments the IP datagram into multiple, self-contained datagrams.



## Fragmentation

- Fragments are re-assembled by the destination host; not by intermediate routers.
- To avoid fragmentation, hosts commonly use path MTU discovery to find the smallest MTU along the path.
- Path MTU discovery involves sending various size datagrams until they do not require fragmentation along the path.
- Most links use  $\text{MTU} \geq 1500$  bytes today.
- Try:  

```
traceroute -F www.uwaterloo.ca 1500 and
traceroute -F www.uwaterloo.ca 1501
```
- (DF=1 set in IP header; routers send "ICMP" error message, which is shown as "!F").
- Bonus: Can you find a destination for which the path MTU < 1500 bytes?

## Fragmentation, Reassembly

*example:*

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

1480 bytes in  
data field  
Subtracted 20 bytes for IP header

Offset (byte addressable)  
= 1480/8

length	ID	fragflag	offset
=4000	=x	=0	=0

*one large datagram becomes  
several smaller datagrams*

length	ID	fragflag	offset
=1500	=x	=1	=0

length	ID	fragflag	offset
=1500	=x	=1	=185


length	ID	fragflag	offset
=1040	=x	=0	=370

## IP Addresses

- IP (Version 4) addresses are 32 bits long
- Every interface has a unique IP address:
  - A computer might have two or more IP addresses
  - A router has many IP addresses
- IP addresses are hierarchical
  - They contain a network ID and a host ID
  - E.g. Apple computers addresses start with: 17....
- IP addresses are assigned statically or dynamically (e.g. DHCP)
- IP (Version 6) addresses are 128 bits long

## IP Addresses

Originally there were 5 classes:



## IP Addresses – Examples


Class “A” address: `www.mit.edu`  
`18.7.22.83`  
 (18 < 128 => Class A)

Class “B” address: `www.toronto.edu`  
`142.150.210.13`  
 (128 < 142 < 128 + 64 => Class B)



## IP Addresses


- **IP address:** 32-bit identifier for host, router interface
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- **IP addresses associated with each interface**





## Subnets

- IP address:
  - subnet part - high order bits
  - host part - low order bits
- *what's a subnet?*
  - device interfaces with same subnet part of IP address
  - can physically reach each other *without intervening router*



## IP Addressing

### Problem:

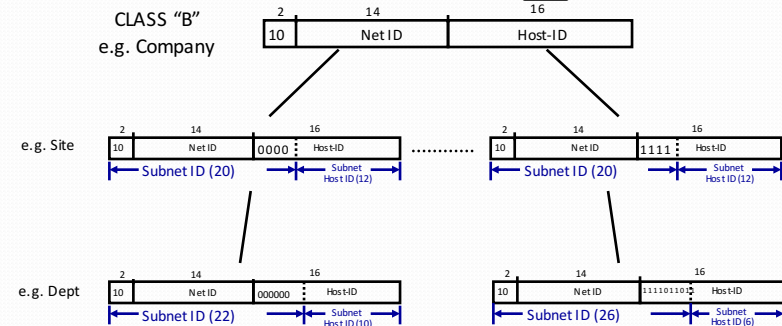
- Address classes were too “rigid”. For most organizations, Class C were too small and Class B too big. Led to inefficient use of address space, and a shortage of addresses.
- Organizations with internal routers needed to have a separate (Class C) network ID for each link.
- And then every other router in the Internet had to know about every network ID in every organization, which led to large address tables.
- Small organizations wanted Class B in case they grew to more than 255 hosts. But there were only about 16,000 Class B network IDs.

## IP Addressing

Two solutions were introduced:

- Subnetting within an organization to subdivide the organization's network ID.
- Classless Inter-Domain Routing (CIDR) in the Internet backbone was introduced in 1993 to provide more efficient and flexible use of IP address space.
- CIDR is also known as “supernetting” because subnetting and CIDR are basically the same idea.

## Subnetting




## Subnetting

- Subnetting is a form of hierarchical routing.
- Subnets are usually represented via an address plus a subnet mask or “netmask”.
- E.g.  


```
yganjali@apps0.cs.toronto.edu > ifconfig eth0
Link encap:Ethernet HWaddr 00:15:17:1C:85:30
inet addr:128.100.3.40 Bcast:128.100.3.255 Mask:ffffff00
UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
```
- Netmask fffffff00: the first 24 bits are the subnet ID, and the last 8 bits are the host ID.
- Can also be represented by a “prefix + length”, e.g. 128.100.3.0/24, or just 128.100.3/24.

## Classless Inter-Domain Routing (CIDR) Addressing

- The IP address space is broken into line segments.
- Each line segment is described by a prefix.
- A prefix is of the form x/y where x indicates the prefix of all addresses in the line segment, and y indicates the length of the segment.
- E.g. The prefix 128.9/16 represents the line segment containing addresses in the range: 128.9.0.0 ... 128.9.255.255.



### Classless Inter-Domain Routing (CIDR) – Addressing



Most specific route = "longest matching prefix"

### Classless Inter-Domain Routing (CIDR) – Addressing

#### Prefix aggregation:

If a service provider serves two organizations with prefixes, it can (sometimes) aggregate them to form a shorter prefix. Other routers can refer to this shorter prefix, and so reduce the size of their address table.

E.g. ISP serves 128.9.14.0/24 and 128.9.15.0/24, it can tell other routers to send it all packets belonging to the prefix 128.9.14.0/23.

#### ISP Choice:

In principle, an organization can keep its prefix if it changes service providers.

## Detour: Map Computer Names to IP addresses


### The Domain Naming System (DNS)

- Names are hierarchical and belong to a domain:
  - e.g. apps0.cs.utoronto.ca
  - Common domain names: .com, .edu, .gov, .org, .net, .ca (or other country-specific domain).
  - Top-level names are assigned by the Internet Corporation for Assigned Names and Numbers (ICANN).
  - A unique name is assigned to each organization.
- DNS Client-Server Model
  - DNS maintains a hierarchical, distributed database of names.
  - Servers are arranged in a hierarchy.
  - Each domain has a “root” server.
  - An application needing an IP address is a DNS client.

## Mapping Computer Names to IP addresses

### The Domain Naming System (DNS)

- A DNS Query
  - Client asks local server.
  - If local server does not have address, it asks the root server of the requested domain.
  - Addresses are cached in case they are requested again.



Example: On CDF machines, try “host [www.eecs.berkeley.edu](http://www.eecs.berkeley.edu)”

## ICMP: internet control message protocol

- used by hosts & routers to communicate network-level information

- error reporting: unreachable host, network, port, protocol
- echo request/reply (used by ping)

- network-layer “above” IP:
  - ICMP msgs carried in IP datagrams

- ICMP message: type, code plus first 8 bytes of IP datagram causing error

Type	Code	description
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

## Traceroute and ICMP

- source sends series of UDP segments to destination

- first set has TTL=1
- second set has TTL=2, etc.
- unlikely port number


- when datagram in  $n$ th set arrives to  $n$ th router:

- router discards datagram and sends source ICMP message (type 11, code 0)
- ICMP message include name of router & IP address

- when ICMP message arrives, source records RTTs

### stopping criteria:

- UDP segment eventually arrives at destination host
- destination returns ICMP “port unreachable” message (type 3, code 3)
- source stops



## An Aside – Error Reporting (ICMP) and traceroute

On CDF machines try: `traceroute www.google.com`  
`traceroute to www.google.com (74.125.159.147), 30 hops max, 40 byte packets`  
 1 `butler.sylab.sandbox (192.168.70.100) 0.103 ms 0.092 ms 0.082 ms`  
 2 `foundry0.cs.toronto.edu (128.100.5.210) 2.146 ms 4.061 ms 5.977 ms`  
 3 `sf-cs1.gw.utoronto.ca (128.100.1.253) 2.184 ms 2.175 ms 2.168 ms`  
 4 `muris-gpb.gw.utoronto.ca (128.100.96.2) 2.146 ms 2.483 ms 3.037 ms`  
 5 `skye2muris-blue.gw.utoronto.ca (128.100.200.210) 7.088 ms 7.207 ms 7.198 ms`  
 6 `muris2skye-yellow.gw.utoronto.ca (128.100.200.217) 3.310 ms 11.325 ms 12.061 ms`  
 7 `ut-hub-utoronto-jf.gtinet.ca (205.211.94.129) 12.681 ms 2.541 ms 2.535 ms`  
 8 `ORION-GTANET-RNE.DIST1-TORO.IPorion.on.ca (66.97.23.57) 3.638 ms 4.391 ms 4.384 ms`  
 9 `BRDR2-TORO-GE2-1.IPorion.on.ca (66.97.16.121) 4.368 ms 4.729 ms 4.844 ms`  
 10 `74.125.51.233 (74.125.51.233) 12.459 ms 12.453 ms 12.808 ms`  
 11 `216.239.47.114 (216.239.47.114) 4.681 ms 4.795 ms 12.661 ms`  
 12 `209.85.250.111 (209.85.250.111) 23.666 ms 23.659 ms 13.226 ms`  
 13 `209.85.242.215 (209.85.242.215) 32.436 ms 32.431 ms 32.913 ms`  
 14 `72.14.232.213 (72.14.232.213) 33.537 ms 72.14.232.215 (72.14.232.215) 33.525 ms 72.14.232.213 (72.14.232.213) 164.315 ms`  
 15 `209.85.254.14 (209.85.254.14) 45.864 ms 209.85.254.10 (209.85.254.10) 42.232 ms 209.85.254.6 (209.85.254.6) 42.346 ms`  
 16 `yi-in-f147.google.com (74.125.159.147) 34.728 ms 34.727 ms 34.713 ms`

## An Aside – Error Reporting (ICMP) and traceroute

### Internet Control Message Protocol

- Used by a router/end-host to report some types of error:
- E.g. Destination Unreachable: packet can't be forwarded to/towards its destination.
- E.g. Time Exceeded: TTL reached zero, or fragment didn't arrive in time. traceroute uses this error to its advantage.
- An ICMP message is an IP datagram, and is sent back to the source of the packet that caused the error.

## Summary

- Shuttling data from one link to another
  - Bits, frames, packets, ...
  - Repeaters/hubs, bridges/switches, routers, ...
- Key ideas in switches
  - Cut-through switching
  - Self learning of the switch table
  - Spanning trees
- Internet Protocol
  - Addresses, subnets, CIDR
  - DNS, Traceroute, ICMP

## Questions?





